

# scrapper\_clean

April 26, 2021

```
[1]: # !pip install webdriver-manager
```

```
[2]: from selenium import webdriver
from selenium.webdriver.common.keys import Keys
from webdriver_manager.chrome import ChromeDriverManager
from selenium.webdriver.chrome.options import Options # for suppressing the
↳ browser
from selenium.webdriver.common.by import By
from selenium.webdriver.support.ui import WebDriverWait
from selenium.webdriver.support import expected_conditions as EC
from bs4 import BeautifulSoup
import re
import pandas as pd
import os
import time
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

## 0.0.1 Scrapping data for Robusta coffee species

```
[3]: cleaned_robusta_df = pd.read_csv('robusta_data_cleaned.csv')
```

### Extract all the rows

```
[4]: url = 'https://database.coffeeinstitute.org/coffees/robusta'

option = webdriver.ChromeOptions()
option.add_argument('headless') # use webdriver without opening the browser
↳ window
driver = webdriver.Chrome(ChromeDriverManager().install(), options=option)

driver.implicitly_wait(30)
driver.get(url)
time.sleep(8)
```

```

===== WebDriver manager =====
Current google-chrome version is 89.0.4389
Get LATEST driver version for 89.0.4389
Driver [/Users/david/.wdm/drivers/chromedriver/mac64/89.0.4389.23/chromedriver]
found in cache

```

```

[5]: robusta_main = BeautifulSoup(driver.page_source, 'lxml')
rows = robusta_main.find_all('tr', class_=['odd', 'even']) # extract odd and
↳ even rows from the table

```

### Extract the url content for each sample

```

[27]: root_url = 'https://database.coffeeinstitute.org'
links = [root_url + row.find('a')['href'] for row in rows] # extract the link
↳ for each row in the database
robusta_df = pd.DataFrame() # initialize df to store all records
print(links[:3])

```

```

['https://database.coffeeinstitute.org/coffee/722152',
'https://database.coffeeinstitute.org/coffee/758792',
'https://database.coffeeinstitute.org/coffee/805984']

```

### Store the information of each sample in a DataFrame

```

[28]: for link in links:
    driver.implicitly_wait(10)
    driver.get(link)
    # Wait until the element with TAG_NAME 'tr' has been loaded
    element = WebDriverWait(driver, 10).until(
        EC.presence_of_element_located((By.TAG_NAME, "tr"))
    )

    info = BeautifulSoup(driver.page_source, 'lxml')
    robusta_dic = {}

    for i in range(1, len(info.find_all('table'))):
        # find all 'th' the headers of the table
        table1 = info.find_all('table')[i].find_all('th')
        table_keys = [x.get_text() for x in table1]
        # find all the data of each header
        table2 = info.find_all('table')[i].find_all('td')
        table_values = [x.get_text() if x.get_text() != '' else np.nan for x in
↳ table2]

        tmp_dic = dict(zip(table_keys, table_values))

```

```

    robusta_dic.update(tmp_dic) # update dictionary with the information
    ↪for robusta species

    df = pd.DataFrame(robusta_dic, index=[0]) # add index=[0] because the
    ↪values are scalars (not in a list)
    robusta_df = pd.concat([robusta_df,df], ignore_index=True, sort=False)
    time.sleep(5)

# quit the driver
# driver.quit()

```

```
[30]: robusta_df.head(3)
```

```

[30]: Country of Origin Number of Bags      Farm Name Bag Weight Lot Number \
0          India          100  Sethuraman Estate      60 kg Lot No 22
1          Mexico          320                n/a      60 kg  1540038
2          India          170  Sethuraman Estate      60 kg          27

```

```

                                In-Country Partner      Mill \
0  NKG Quality Service (a division of Bernhard Ro...  Kaapi Royale
1  NKG Quality Service (a division of Bernhard Ro...  AMSA - ECOM
2  NKG Quality Service (a division of Bernhard Ro...  Kaapi Royale

```

```

Harvest Year      ICO Number      Grading Date ... Moisture \
0      2020  14/1148/2020/11  August 20th, 2020 ...    12 %
1      2019   016-2222-0409  November 3rd, 2020 ...    10 %
2      2020  14/1148/2020/12  August 20th, 2020 ...    12 %

```

```

Color Category One Defects Category Two Defects Quakers \
0      Green      0 full defects      0 full defects      0 NaN
1  Yellow-Green      1 full defects      7 full defects      3 NaN
2      Green      0 full defects      0 full defects      0 NaN

```

```

Expiration      Certification Body \
0  August 20th, 2021  NKG Quality Service (a division of Bernhard Ro...
1  November 3rd, 2021  NKG Quality Service (a division of Bernhard Ro...
2  August 20th, 2021  NKG Quality Service (a division of Bernhard Ro...

```

```

Certification Address      Certification Contact
0  Bahnhofstrasse 22 6300 Zug, Switzerland  Gloria Pedroza - +41417287296
1  Bahnhofstrasse 22 6300 Zug, Switzerland  Gloria Pedroza - +41417287296
2  Bahnhofstrasse 22 6300 Zug, Switzerland  Gloria Pedroza - +41417287296

```

```
[3 rows x 40 columns]
```

## 0.0.2 Scrapping data for Arabica coffee species

```
[31]: cleaned_arabica_df = pd.read_csv('arabica_data_cleaned.csv')
```

Extract all the rows from different pages

```
[32]: url = 'https://database.coffeeinstitute.org/coffees/arabica'

option = webdriver.ChromeOptions()
option.add_argument('headless')
driver = webdriver.Chrome(ChromeDriverManager().install(), options=option)

driver.implicitly_wait(30)
driver.get(url)
time.sleep(8)

arabica_main = BeautifulSoup(driver.page_source, 'lxml')
rows = arabica_main.find_all('tr', class_='odd even') # extract odd and
↳ even rows from the table

while len(arabica_main.find_all('a', class_='paginate_button next disabled')) <
↳ 1:
    page_buttons = driver.find_elements_by_class_name('paginate_button')
    page_buttons[-1].click() # click next
    time.sleep(6)
    arabica_main = BeautifulSoup(driver.page_source, 'lxml') # get arabica_main
↳ for next page
    rows += arabica_main.find_all('tr', class_='odd even') # add new rows
```

===== WebDriver manager =====

Current google-chrome version is 89.0.4389

Get LATEST driver version for 89.0.4389

Driver [/Users/david/.wdm/drivers/chromedriver/mac64/89.0.4389.23/chromedriver]  
found in cache

Extract the url content for each sample

```
[33]: root_url = 'https://database.coffeeinstitute.org'
links = [root_url + row.find('a')['href'] for row in rows] # extract the link
↳ for each row in the database
arabica_df = pd.DataFrame() # initialize df to store all records
print(links[-3:])
```

```
['https://database.coffeeinstitute.org/coffee/813284',
'https://database.coffeeinstitute.org/coffee/564165',
```

```
'https://database.coffeeinstitute.org/coffee/799551']
```

Store the information for each sample in a DataFrame

```
[34]: for link in links:
        driver.implicitly_wait(10)
        driver.get(link)
        # Wait until the element with TAG_NAME 'tr' has been loaded
        element = WebDriverWait(driver, 10).until(
            EC.presence_of_element_located((By.TAG_NAME, "tr"))
        )

        info = BeautifulSoup(driver.page_source, 'lxml')
        arabica_dic = {}

        for i in range(1, len(info.find_all('table'))):
            # find all 'th' the headers of the table
            table1 = info.find_all('table')[i].find_all('th')
            table_keys = [x.get_text() for x in table1]
            # find all the data of each header
            table2 = info.find_all('table')[i].find_all('td')
            table_values = [x.get_text() if x.get_text() != '' else np.nan for x in
→table2]

            tmp_dic = dict(zip(table_keys, table_values))
            arabica_dic.update(tmp_dic) # update dictionary with the information
→for arabica species

            df = pd.DataFrame(arabica_dic, index=[0]) # add index=[0] because the
→values are scalars (not in a list)
            arabica_df = pd.concat([arabica_df, df], ignore_index=True, sort=False)
            time.sleep(4)

        # quit the driver
        # driver.quit()
```

```
[39]: arabica_df.head(3)
```

```
[39]: Country of Origin Number of Bags Farm Name Bag Weight Lot Number \
0 Taiwan 3 60 kg 202004
1 Kenya 320 N/A 69 kg 1542247
2 Ethiopia 37 Honey Drip 20 kg 1

In-Country Partner Mill \
0 NKG Quality Service (a division of Bernhard Ro...
1 NKG Quality Service (a division of Bernhard Ro... Tylor Winch
2 Japan Coffee Exchange non
```

	Harvest Year	ICO Number	Grading Date	...	Moisture	Color	\
0	2020	NaN	September 1st, 2020	...	11 %	None	
1	2019	037-1673-5107	October 30th, 2020	...	11 %	Green	
2	2020	NaN	July 14th, 2020	...	11 %	Green	

  

	Category One Defects	Category Two Defects	Quakers	Expiration	\
0	0 full defects	0 full defects	0 NaN	September 1st, 2021	
1	0 full defects	1 full defects	0 NaN	October 30th, 2021	
2	0 full defects	4 full defects	3 NaN	July 14th, 2021	

  

	Certification Body	\
0	NKG Quality Service (a division of Bernhard Ro...	
1	NKG Quality Service (a division of Bernhard Ro...	
2	Japan Coffee Exchange	

  

	Certification Address	\
0	Bahnhofstrasse 22 6300 Zug, Switzerland	
1	Bahnhofstrasse 22 6300 Zug, Switzerland	
2	413-0002 - 1173-58 Izu...	

  

	Certification Contact
0	Gloria Pedroza - +41417287296
1	Gloria Pedroza - +41417287296
2	Koju Matsuzawa - +81(0)9085642901

[3 rows x 40 columns]

### 0.0.3 Cleaning

```
[343]: # Drop empty column (after Quaker)
arabica_df.drop('', axis=1, inplace=True)
robusta_df.drop('', axis=1, inplace=True)

[370]: # Change the names of the columns to match the format of the cleaned datasets
new_names = []
for col in arabica_df.columns:
    new_names.append(col.replace(' ', '.').replace('-', '.'))

arabica_df.columns = new_names
robusta_df.columns = new_names

[374]: # Store the raw data
arabica_df.to_csv('arabica_raw.csv', index=False)
robusta_df.to_csv('robusta_raw.csv', index=False)
```

```
[390]: # arabica_df = pd.read_csv('arabica_raw.csv')
# robusta_df = pd.read_csv('robusta_raw.csv')
```

```
[391]: # Region, Owner, Company, Farm.Name and Mill to lower
cols_lower = ['Region', 'Owner', 'Company', 'Farm.Name', 'Mill']
for col in cols_lower:
    arabica_df.loc[:,col] = arabica_df.loc[:,col].str.lower()
    robusta_df.loc[:,col] = robusta_df.loc[:,col].str.lower()
```

```
[392]: # Add column Species and Unnamed: 0 which keeps an index
arabica_df['Species'] = 'Arabica'
robusta_df['Species'] = 'Robusta'
arabica_df.index += 1313
robusta_df.index += 29
arabica_df.index.rename('Unnamed: 0', inplace=True)
robusta_df.index.rename('Unnamed: 0', inplace=True)
arabica_df.reset_index(inplace=True)
robusta_df.reset_index(inplace=True)
arabica_df.head(3)
```

```
[392]: Unnamed: 0 Country.of.Origin Number.of.Bags Farm.Name Bag.Weight \
0      1313      Taiwan      3      60 kg
1      1314      Kenya     320      NaN      69 kg
2      1315      Ethiopia     37  honey drip      20 kg

Lot.Number In.Country.Partner Mill \
0      202004 NKG Quality Service (a division of Bernhard Ro...
1      1542247 NKG Quality Service (a division of Bernhard Ro... tylor winch
2      1      Japan Coffee Exchange non

Harvest.Year ICO.Number ... Moisture Color Category.One.Defects \
0      2020      NaN ... 11 % None 0 full defects
1      2019 037-1673-5107 ... 11 % Green 0 full defects
2      2020      NaN ... 11 % Green 0 full defects

Category.Two.Defects Quakers Expiration \
0      0 full defects 0 September 1st, 2021
1      1 full defects 0 October 30th, 2021
2      4 full defects 3 July 14th, 2021

Certification.Body \
0 NKG Quality Service (a division of Bernhard Ro...
1 NKG Quality Service (a division of Bernhard Ro...
2      Japan Coffee Exchange

Certification.Address \
0      Bahnhofstrasse 22 6300 Zug, Switzerland
```

```

1      Bahnhofstrasse 22 6300 Zug, Switzerland
2  413-0002      -      1173-58 Izu...

```

```

      Certification.Contact  Species
0      Gloria Pedroza - +41417287296  Arabica
1      Gloria Pedroza - +41417287296  Arabica
2      Koju Matsuzawa - +81(0)9085642901  Arabica

```

[3 rows x 41 columns]

Create **altitude\_low\_meters**, **altitude\_high\_meters**, **altitude\_mean\_meters** Some of the values are very difficult to fix using regular expressions or an automated command, but besides 2 particular cases, we can use str methods to clean the column of altitude and get the values of the previously cleaned datasets to join them together.

```

[461]: # 2 special cases to be cleaned
arabica_df.loc[4, 'Altitude'] = '1500-2100'
arabica_df.loc[90, 'Altitude'] = '1100'

# use regex to clean data and split by - into two columns (low and high
↳altitude)
df = arabica_df.Altitude.str.strip().str.replace(',|\.\d*|\\+', '').str.
↳split('-', expand=True)
df.columns = ['altitude_low_meters', 'altitude_high_meters']

# make the None values in high column equal to the low column
df.altitude_high_meters[pd.isna(df.altitude_high_meters)] = df.
↳altitude_low_meters[pd.isna(df.altitude_high_meters)]
df.altitude_low_meters[df.altitude_low_meters==''] = df.altitude_high_meters[df.
↳altitude_low_meters=='']
df['altitude_mean_meters'] = (df.altitude_low_meters.astype(int) + df.
↳altitude_high_meters.astype(int))/2

arabica_df = pd.concat([arabica_df, df], axis=1)

```

```

[468]: # No need to preprocess the Robusta Altitude column, just match the format of
↳Arabica
robusta_df.Altitude = robusta_df.Altitude.astype(int)
robusta_df['altitude_low_meters'] = robusta_df['Altitude']
robusta_df['altitude_high_meters'] = robusta_df['Altitude']
robusta_df['altitude_mean_meters'] = robusta_df['Altitude']

```

Remove extra columns in all datasets



```
[493]: # drop columns not present in previously cleaned data
drop_cols = ['Status', 'Overall', 'Defects', 'Certification.Address',
             'Certification.Contact']
arabica_df.drop(drop_cols, axis=1, inplace=True)
robusta_df.drop(drop_cols, axis=1, inplace=True)

# drop columns not present in new scrapped data
drop_cols2 = ['Owner.1', 'Cupper.Points', 'unit_of_measurement', 'Certification.
→Address',
             'Certification.Contact']
cleaned_arabica_df.drop(drop_cols2, axis=1, inplace=True)
cleaned_robusta_df.drop(drop_cols2, axis=1, inplace=True)
```

To be able to do the final join of all the datasets, it is necessary to rename the columns of the previously cleaned robusta data. The following columns will be renamed to match with the current convention for robusta and arabica.

- Fragrance...Aroma -> Aroma
- Salt...Acid -> Acidity
- Bitter...Sweet -> Sweetness
- Mouthfeel -> Body
- Uniform.Cup -> Uniformity

```
[530]: # rename columns in old robusta to match format of all datasets
robusta_dict = {'Fragrance...Aroma': 'Aroma',
               'Salt...Acid': 'Acidity', 'Mouthfeel': 'Body',
               'Bitter...Sweet': 'Sweetness',
               'Uniform.Cup': 'Uniformity'}

cleaned_robusta_df.rename(columns=robusta_dict, inplace=True)
```

#### 0.0.4 Join all dataframes

```
[536]: # concatenate all dataframes by default outter, but all columns match with each_
→other
full_df = pd.concat([cleaned_arabica_df, arabica_df,
                    cleaned_robusta_df, robusta_df],
                    ignore_index=True, sort=False)
```

```
[575]: # Convert Moisture to float
for i, val in enumerate(full_df.Moisture):
    try:
        float(val)
    except:
        full_df.loc[i, 'Moisture'] = float(val.replace('%', '').strip())/100
```

```
full_df.Moisture = full_df.Moisture.astype(float)
```

```
[572]: # Conver Category.One.Defects to int
for i,val in enumerate(full_df['Category.One.Defects']):
    try:
        int(val)
    except:
        full_df.loc[i,'Category.One.Defects'] = int(val.replace('full_
↳defects','').strip())
```

### Checking numeric variables

```
[577]: full_df.describe()
```

```
[577]:
```

	Unnamed: 0	Number.of.Bags	Aroma	Flavor	Aftertaste \
count	1477.000000	1477.000000	1477.000000	1477.000000	1477.000000
mean	704.918754	161.854435	7.574821	7.531774	7.407928
std	425.286918	135.952825	0.372236	0.393496	0.397649
min	1.000000	0.000000	0.000000	0.000000	0.000000
25%	335.000000	16.000000	7.420000	7.330000	7.250000
50%	704.000000	200.000000	7.580000	7.580000	7.420000
75%	1073.000000	275.000000	7.750000	7.750000	7.670000
max	1443.000000	1280.000000	8.750000	8.830000	8.670000

  

	Acidity	Body	Balance	Uniformity	Clean.Cup \
count	1477.000000	1477.000000	1477.000000	1477.000000	1477.000000
mean	7.545166	7.527393	7.526242	9.841273	9.843290
std	0.374695	0.363753	0.400143	0.549402	0.740145
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	7.330000	7.330000	7.330000	10.000000	10.000000
50%	7.580000	7.500000	7.500000	10.000000	10.000000
75%	7.750000	7.750000	7.750000	10.000000	10.000000
max	8.750000	8.580000	8.750000	10.000000	10.000000

  

	Sweetness	Total.Cup.Points	Moisture	Category.One.Defects \
count	1477.000000	1477.000000	1477.000000	1477.000000
mean	9.856994	82.161381	0.090150	0.444144
std	0.612304	3.447233	0.046726	2.433762
min	0.000000	0.000000	0.000000	0.000000
25%	10.000000	81.170000	0.100000	0.000000
50%	10.000000	82.580000	0.110000	0.000000
75%	10.000000	83.670000	0.120000	0.000000
max	10.000000	90.580000	0.280000	63.000000

  

	Quakers	altitude_mean_meters
count	1476.000000	1247.000000

mean	0.266938	1732.473837
std	1.076978	8176.221765
min	0.000000	1.000000
25%	0.000000	1100.000000
50%	0.000000	1325.000000
75%	0.000000	1600.000000
max	18.000000	190164.000000

Above we observe that `altitude_mean_meters` has an unrealistic max value which can greatly affect distance based clustering methods. For this reason, the model will focus on records below 5000, assuming that this is an error in the scale or imputation.

```
[588]: full_df = full_df[full_df['altitude_mean_meters']<5000].copy()
```

```
[1009]: full_df.to_csv('arabica_robusta_cleaned.csv', index=False)
```

```
[ ]:
```